

09RecommendationSystem

March 18, 2019

1 Recommendation

1. 推薦系統的種類

- collabrative(協同推薦)
- content-based(內容推薦)

from [quora](#)

1.1 IMPORT & DATA

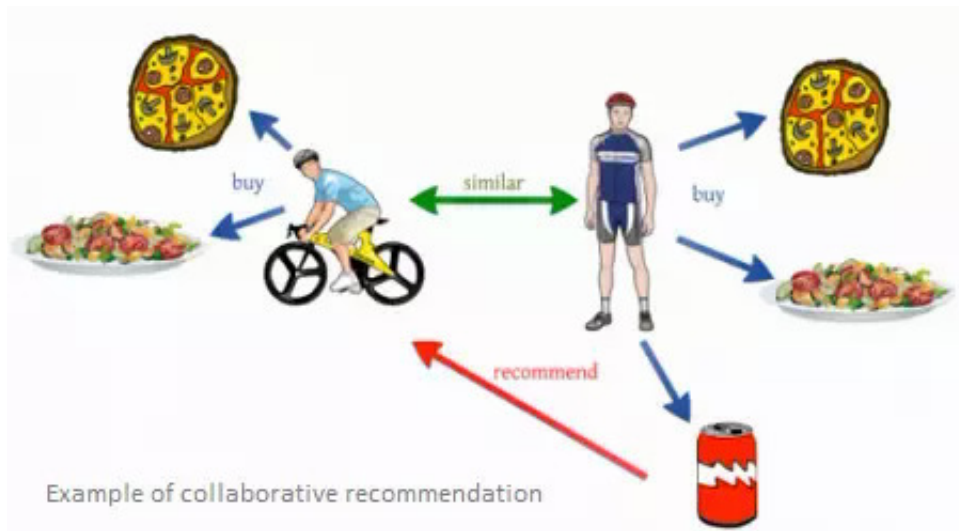
```
In [1]: import numpy as np
import scipy
import pandas as pd
import math
import random
import sklearn
from nltk.corpus import stopwords
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
from scipy.sparse.linalg import svds
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings('ignore')

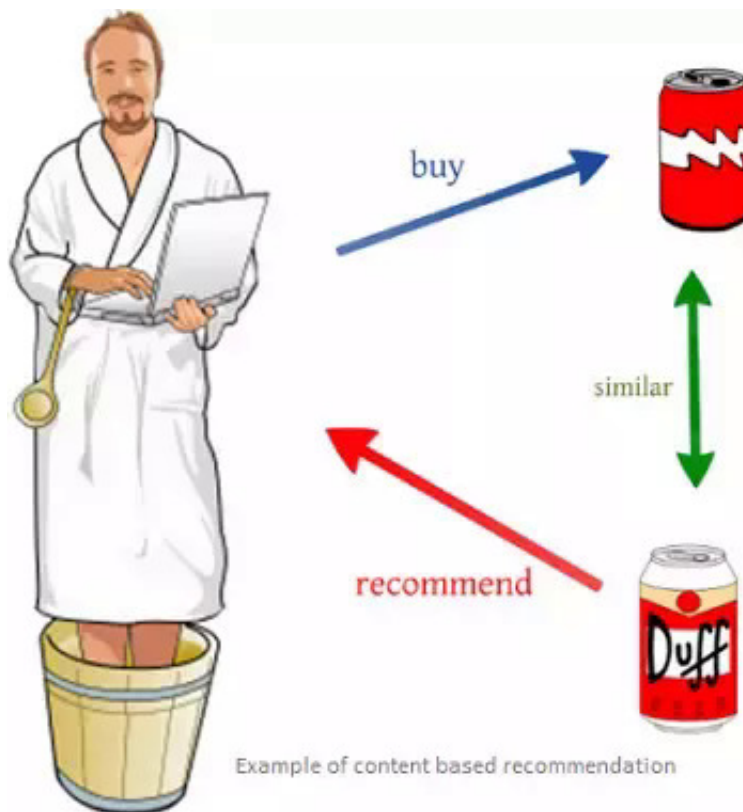
In [2]: df_articles = pd.read_csv('shared_articles.csv')
df_articles = df_articles[df_articles['eventType'] == 'CONTENT SHARED']
df_articles.head(5)

Out[2]:
```

	timestamp	eventType	contentId	authorPersonId	\
1	1459193988	CONTENT SHARED	-4110354420726924665	4340306774493623681	
2	1459194146	CONTENT SHARED	-7292285110016212249	4340306774493623681	
3	1459194474	CONTENT SHARED	-6151852268067518688	3891637997717104548	



collabrative



content-based

```

4 1459194497 CONTENT SHARED 2448026894306402386 4340306774493623681
5 1459194522 CONTENT SHARED -2826566343807132236 4340306774493623681

```

```

      authorSessionId authorUserAgent authorRegion authorCountry contentType \
1 8940341205206233829           NaN           NaN           NaN        HTML
2 8940341205206233829           NaN           NaN           NaN        HTML
3 -1457532940883382585           NaN           NaN           NaN        HTML
4 8940341205206233829           NaN           NaN           NaN        HTML
5 8940341205206233829           NaN           NaN           NaN        HTML

```

```

                                url \
1 http://www.nytimes.com/2016/03/28/business/dea...
2 http://cointelegraph.com/news/bitcoin-future-w...
3 https://cloudplatform.googleblog.com/2016/03/G...
4 https://bitcoinmagazine.com/articles/ibm-wants...
5 http://www.coindesk.com/ieee-blockchain-oxford...

```

```

                                title \
1 Ethereum, a Virtual Currency, Enables Transact...
2 Bitcoin Future: When GBPcoin of Branson Wins O...
3           Google Data Center 360ř Tour
4 IBM Wants to "Evolve the Internet" With Blockc...
5 IEEE to Talk Blockchain at Cloud Computing Oxf...

```

```

                                text lang
1 All of this work is still very early. The firs...   en
2 The alarm clock wakes me at 8:00 with stream o...   en
3 We're excited to share the Google Data Center ...   en
4 The Aite Group projects the blockchain market ...   en
5 One of the largest and oldest organizations fo...   en

```

```

In [3]: df_interactions = pd.read_csv('users_interactions.csv')
df_interactions.head(10)

```

```

Out[3]:   timestamp eventType      contentId      personId \
0 1465413032      VIEW -3499919498720038879 -8845298781299428018
1 1465412560      VIEW  8890720798209849691 -1032019229384696495
2 1465416190      VIEW  310515487419366995 -1130272294246983140
3 1465413895  FOLLOW  310515487419366995   344280948527967603
4 1465412290      VIEW -7820640624231356730 -445337111692715325
5 1465413742      VIEW  310515487419366995 -8763398617720485024
6 1465415950      VIEW -8864073373672512525 3609194402293569455
7 1465415066      VIEW -1492913151930215984 4254153380739593270
8 1465413762      VIEW  310515487419366995 344280948527967603
9 1465413771      VIEW  3064370296170038610 3609194402293569455

```

```

                                sessionId      userAgent \
0 1264196770339959068                                     NaN

```

```

1 3621737643587579081 Mozilla/5.0 (Macintosh; Intel Mac OS X 10_11_2...
2 2631864456530402479 NaN
3 -3167637573980064150 NaN
4 5611481178424124714 NaN
5 1395789369402380392 Mozilla/5.0 (Windows NT 10.0; WOW64) AppleWebK...
6 1143207167886864524 NaN
7 8743229464706506141 Mozilla/5.0 (X11; Linux x86_64) AppleWebKit/53...
8 -3167637573980064150 NaN
9 1143207167886864524 NaN

```

```

      userRegion userCountry
0          NaN          NaN
1           NY           US
2          NaN          NaN
3          NaN          NaN
4          NaN          NaN
5           MG           BR
6          NaN          NaN
7           SP           BR
8          NaN          NaN
9          NaN          NaN

```

1.2 Preprocessing

```

In [4]: set(df_interactions['eventType'])

Out[4]: {'BOOKMARK', 'COMMENT CREATED', 'FOLLOW', 'LIKE', 'VIEW'}

In [5]: event_type_strength = {
        'VIEW': 1.0,
        'LIKE': 2.0,
        'BOOKMARK': 2.5,
        'FOLLOW': 3.0,
        'COMMENT CREATED': 4.0,
    }

# 請將 eventType 按照 event_type_strength 進行評分
#=====your works starts=====#
df_interactions['eventStrength'] =
#=====your works ends=====#

print("平均互動分數:", np.average(df_interactions['eventStrength']))
# 平均互動分數: 1.2362885828078327

```

平均互動分數: 1.2362885828078327

在推薦系統中，有一個很常見的問題，稱為 *cold-start*。因為，很多使用者並沒有真正的根產品產生任何互動，所以並沒有辦法從資料及當中了解到他們偏好。因此，這邊我們將少於五個 interactions 的 user 刪掉。

```

In [6]: # 計算出每個使用者有對幾項不同的商品進行互動 (hint: 以 ['personId', 'contentId'] 進行 groupby)
# 篩選掉互動商品數小於五次的使用者 (hint: return list)
#=====your works starts=====#
df_users_interactions_count =
users_with_enough_interactions =
#=====your works ends=====#

print("平均互動次數: ", np.average(df_users_interactions_count), "次")
print("使用者個數: ", len(df_users_interactions_count))
print("互動大於 5 次使用者個數", len(users_with_enough_interactions))
# 平均互動次數: 21.482849604221634 次
# 使用者個數: 1895
# 互動大於 5 次使用者個數 1140

```

平均互動次數: 21.482849604221634 次
 使用者個數: 1895
 互動大於 5 次使用者個數 1140

```

In [7]: # 找出 df_interactions 中 personId 在 users_with_enough_interactions 當中的 row
#=====your works starts=====#
df_interactions_from_selected_users =
#=====your works ends=====#

print('總互動比數:', len(df_interactions))
print('互動次數大於五用戶總互動比數:', len(df_interactions_from_selected_users))
# 總互動比數: 72312
# 互動次數大於五用戶總互動比數: 69868

```

總互動比數: 72312
 互動次數大於五用戶總互動比數: 69868

```

In [8]: def smooth_user_preference(x):
# 請先 +1 再取 log, 以平滑互動分數
#=====your works starts=====#
logged =
#=====your works ends=====#
return logged

print(smooth_user_preference(1))
print(smooth_user_preference(2))
print(smooth_user_preference(3))
# 0.6931471805599453
# 1.0986122886681098
# 1.3862943611198906

```

0.6931471805599453
 1.0986122886681098

1.3862943611198906

```
In [9]: # 計算使用者對於文章喜愛的加總
# 並透過上面的 smooth_user_preference 使其平滑
#=====your works starts=====#
df_interactions_full_no_smooth =
df_interactions_full =
#=====your works ends=====#

print('平均喜好分數 (未平滑)', np.average(df_interactions_full_no_smooth))
print('平均喜好分數 (平滑)', np.average(df_interactions_full['eventStrength']))
# 平均喜好分數 (未平滑) 2.214954226972843
# 平均喜好分數 (平滑) 1.015265936675581
```

平均喜好分數 (未平滑) 2.214954226972843

平均喜好分數 (平滑) 1.015265936675581

1.3 TRAIN_TEST_SPLIT

```
In [10]: # # 請使用 train_test_split 切分 df_interactions_full
# # 1. stratify=df_interactions_full['personId'] #stratify 可以按照 y 的比例進行切分
# # 2. test_size=0.2
# # 3. random_state=1212
# #=====your works starts=====#
# df_interactions_train, df_interactions_test =
# #=====your works ends=====#

# print('len(df_interactions_train):', len(df_interactions_train))
# print('len(df_interactions_test):', len(df_interactions_test))
# # len(df_interactions_train): 31284
# # len(df_interactions_test): 7822
```

1.4 評價

1.4.1 基礎知識

col	Retrieved	Non Retrieved
Relevant	True Positive(TP)	False Negative(TN)
Irrelevant	False Postive(FP)	True Negative(TN)

1.4.2 解釋

1. Precision at K

$$Precision = \frac{|Relevant \cap Retrieved|}{|Retrieved|} = \frac{|TP|}{|TP| + |FP|}$$

2. Recall at K

$$Recall = \frac{|Relevant \cap Retrieved|}{|Relevant|} = \frac{|TP|}{|TP| + |TN|}$$

3. F measure

1. 算術平均數

$$F = \frac{P + R}{2} \quad \text{where } P = Precision, R = Recall$$

2. 幾何平均數

$$F = \frac{1}{\alpha \frac{1}{P} + (1 - \alpha) \frac{1}{R}} = \frac{(\beta^2 + 1)PR}{\beta^2 P + R} \quad \text{where } \beta^2 = \frac{1 - \alpha}{\alpha}$$

3. F_β or F_1

$$F_1 = F_\beta = \frac{1}{0.5 \frac{1}{P} + 0.5 \frac{1}{R}} = \frac{PR}{0.5P + 0.5R} = \frac{2PR}{P + R}$$

4. R-Precision

5. NDCG

6. MAP

1.4.3 注意事項

1. 注意與 Accuracy 的區別

$$\frac{|TP| + |TF|}{|TP| + |TN| + |FP| + |FN|}$$

2. 參考資訊

1.5 Evaluation

```
In [11]: # predict result
item_ids = df_articles['contentId'].tolist()
item_ids_mapping = dict([(idx, contentId) for idx, contentId in enumerate(item_ids)])
item_ids_set = set(item_ids)
df_user_preference = df_interactions_full.groupby('personId')['contentId'].apply(set)
```

```
In [12]: # ground truth
# 找出每一個 personId 曾經互動過「不重複」的文章
#=====your works starts=====#
df_answer =
#=====your works ends=====#

df_answer.head(5)
# personId
# -9223121837663643404    {5211673327552264703, -5002383425685129595, -7...
# -9212075797126931087    {-1995591062742965408, 6852597772196653540, -9...
# -9207251133131336884    {-4029704725707465084, -1297580205670251233, -...
# -9199575329909162940    {5293701842202310496, -5002383425685129595, 54...
# -9196668942822132778    {-721732705314803549, -8813724423497152538, -8...
# Name: contentId, dtype: object
```

```
Out[12]: personId
-9223121837663643404    {5211673327552264703, -5002383425685129595, -7...
-9212075797126931087    {-1995591062742965408, 6852597772196653540, -9...
-9207251133131336884    {-4029704725707465084, -1297580205670251233, -...
-9199575329909162940    {5293701842202310496, -5002383425685129595, 54...
-9196668942822132778    {-721732705314803549, -8813724423497152538, -8...
Name: contentId, dtype: object
```

1.6 Popularity model (Base Line)

```
In [13]: # 以 contentId 進行 groupby · 按照每篇文章總分數進行排序
#=====your works starts=====#
df_item_popularity =
#=====your works ends=====#

df_item_popularity.head(5).to_dict(orient='record')
# [{'contentId': -4.029704725707465e+18, 'eventStrength': 213.30481497288199},
#  {'contentId': -6.783772548752092e+18, 'eventStrength': 162.03158006500846},
#  {'contentId': -1.3313934239753886e+17, 'eventStrength': 158.05458586966674},
#  {'contentId': -8.208801367848628e+18, 'eventStrength': 136.62458307425328},
#  {'contentId': -6.843047699859122e+18, 'eventStrength': 134.34939619163308}]
```

```
Out[13]: [{'contentId': -4.029704725707465e+18, 'eventStrength': 213.30481497288199},
{'contentId': -6.783772548752092e+18, 'eventStrength': 162.03158006500846},
{'contentId': -1.3313934239753886e+17, 'eventStrength': 158.05458586966674},
{'contentId': -8.208801367848628e+18, 'eventStrength': 136.62458307425328},
{'contentId': -6.843047699859122e+18, 'eventStrength': 134.34939619163308}]
```

```
In [14]: def popularity_recommend(user_id):
# 直接回傳分數加總最高的十篇文章
#=====your works starts=====#
recommend =
#=====your works ends=====#
return recommend
```



```
# 透過 apply function 使用 popularity_recommend 到 df_interactions_full["personId"] 的
df_user_preference['popularity_recommend'] = df_user_preference["personId"].apply(popularity_recommend)
df_user_preference.head(5)
```

```
Out [14]:
```

	personId	contentId \
0	-9223121837663643404	{5211673327552264703, -5002383425685129595, -7...
1	-9212075797126931087	{-1995591062742965408, 6852597772196653540, -9...
2	-9207251133131336884	{-9216926795620865886, -4029704725707465084, -...
3	-9199575329909162940	{5293701842202310496, -5002383425685129595, 54...
4	-9196668942822132778	{-721732705314803549, -8813724423497152538, -8...


```
popularity_recommend
0 [-4029704725707465084, -6783772548752091658, -...
1 [-4029704725707465084, -6783772548752091658, -...
2 [-4029704725707465084, -6783772548752091658, -...
3 [-4029704725707465084, -6783772548752091658, -...
4 [-4029704725707465084, -6783772548752091658, -...
```

```
In [15]: def precision_at_k(row, k=10):
# 計算每一個 row 的 precision_at_k
#=====your works starts=====#
precision =
#=====your works ends=====#
return precision

evaluation_result = df_user_preference.apply(precision_at_k, axis=1)
print("Average Precision At K:", np.average(evaluation_result))
```

Average Precision At K: 0.13342105263157894

1.7 Content-Based Filtering model

```
In [16]: stopwords_list = stopwords.words('english') + stopwords.words('portuguese')
vectorizer = TfidfVectorizer(analyzer='word', ngram_range=(1, 2), min_df=0.003, max_d...
```

```
In [17]: # 計算 df_articles['title'] + " " + df_articles['text'] 中每個 row 的 tfidf_matrix
#=====your works starts=====#
tfidf_matrix =
#=====your works ends=====#

np.sum(tfidf_matrix[:5].toarray(), axis=1)
# array([ 9.50309706,  9.14139363,  7.07473481, 10.43412109,  7.64140829])
```

```
Out [17]: array([ 9.50309706,  9.14139363,  7.07473481, 10.43412109,  7.64140829])
```

```
In [18]: def get_user_vector(like_content_list):
idxs = [item_ids.index(i) for i in like_content_list if i in item_ids_set]
```

```

if len(idxs) == 0:
    # 初始劃一條全部為零，與 tfidf_matrix 中每一條向量等長的 np.array()
    #=====your works starts=====#
    average_vector =
    #=====your works ends=====#
else:
    # 使用 idxs 找出 tfidf_matrix 中的對應向量
    # 並 element-wise 的計算每一條向量中每個元素的平均值 (axis=0)
    #=====your works starts=====#
    tfidf_vectors =
    average_vector =
    #=====your works ends=====#
return average_vector

df_user_preference['preference_vector'] = df_user_preference['contentId'].apply(get_u
df_user_preference['preference_vector'].head(5)
# 0    [0.0022163282029446594, 0.0030194389616064613,...
# 1    [0.0, 0.012492323019192434, 0.0, 0.0, 0.003230...
# 2    [0.013699793673353864, 0.002006773900308267, 0...
# 3    [0.0, 0.006253150796304994, 0.0, 0.0, 0.003906...
# 4    [0.0, 0.0, 0.0, 0.0, 0.017006668208970826, 0.0...
# Name: preference_vector, dtype: object

Out[18]: 0    [0.0022163282029446594, 0.0030194389616064613,...
1    [0.0, 0.012492323019192434, 0.0, 0.0, 0.003230...
2    [0.013699793673353864, 0.002006773900308267, 0...
3    [0.0, 0.006253150796304994, 0.0, 0.0, 0.003906...
4    [0.0, 0.0, 0.0, 0.0, 0.017006668208970826, 0.0...
Name: preference_vector, dtype: object

In [19]: user_preference_vector = np.hstack(df_user_preference['preference_vector'].values).res
# 使用 cosine_similarity 去計算每一個 preference_vector 與每一篇待選文章的 cosine simil
#=====your works starts=====#
similarity_metric =
#=====your works ends=====#
# 請特別注意：
# 每一個 row 是使用者對每一篇文章的 similarity(preference)
# 所以接下來要篩出，每一個 row 當中 similarity 最高分的 10 篇文章

similarity_metric[:5, :5]
# array([[0.119037 , 0.11127776, 0.23397873, 0.16428804, 0.1500561 ],
#        [0.03351941, 0.03650116, 0.13718512, 0.04648663, 0.04216605],
#        [0.04209047, 0.04084319, 0.04926266, 0.08144102, 0.02540297],
#        [0.08642438, 0.09372568, 0.10180722, 0.1332371 , 0.08941768],
#        [0.04121156, 0.01011304, 0.01477526, 0.04543164, 0.03242129]])

Out[19]: array([[0.119037 , 0.11127776, 0.23397873, 0.16428804, 0.1500561 ],
               [0.03351941, 0.03650116, 0.13718512, 0.04648663, 0.04216605],

```

```
[0.04209047, 0.04084319, 0.04926266, 0.08144102, 0.02540297],
[0.08642438, 0.09372568, 0.10180722, 0.1332371 , 0.08941768],
[0.04121156, 0.01011304, 0.01477526, 0.04543164, 0.03242129]])
```

```
In [20]: # 使用 np.argsort 將每一個 row 的 similarity 進行排序，然後倒過來排續，篩出前 10 個
#=====your works starts=====#
top_10_content_idx =
#=====your works ends=====#
# 請注意這邊的產出代表的是每一篇文章的在 tfidf_matrix 的 idx 位置
# 必須與 contentId 區別
```

```
top_10_content_idx
# array([[ 650, 1032, 1643, ..., 237, 1034, 3018],
#        [ 977, 1023, 1601, ..., 1548, 1175, 1769],
#        [1671, 1593, 1795, ..., 2477, 1117, 1520],
#        ...,
#        [1622, 1845, 3021, ..., 974, 1035, 1607],
#        [1185, 1636, 1116, ..., 2309, 2357, 2616],
#        [2664, 2781, 659, ..., 2634, 1552, 3018]], dtype=int64)
```

```
Out[20]: array([[ 650, 1032, 1643, ..., 237, 1034, 3018],
[ 977, 1023, 1601, ..., 1548, 1175, 1769],
[1671, 1593, 1795, ..., 2477, 1117, 1520],
...,
[1622, 1845, 3021, ..., 974, 1035, 1607],
[1185, 1636, 1116, ..., 2309, 2357, 2616],
[2664, 2781, 659, ..., 2634, 1552, 3018]], dtype=int64)
```

```
In [21]: # 將 tfidf_matrix 的 idx 轉換成 contentId
#=====your works starts=====#
top_10_contentId =
#=====your works ends=====#

top_10_contentId[:5, :3]
# array([[ 8596997246990922861, 2858969450431709251, -4541461982704074404],
#        [-1995591062742965408, 6852597772196653540, -969155230116728853],
#        [-1297580205670251233, -9216926795620865886, -4434534460030275781],
#        [-1755875383603052680, 5293701842202310496, 5037403311832115000],
#        [ 9175693555063886126, 7013665235990336340, -2069509552243850466]],
#        dtype=int64)
```

```
Out[21]: array([[ 8596997246990922861, 2858969450431709251, -4541461982704074404],
[-1995591062742965408, 6852597772196653540, -969155230116728853],
[-1297580205670251233, -9216926795620865886, -4434534460030275781],
[-1755875383603052680, 5293701842202310496, 5037403311832115000],
[ 9175693555063886126, 7013665235990336340, -2069509552243850466]],
dtype=int64)
```

```
In [22]: df_user_preference['content_based_recommended'] = list(top_10_contentId)
```

```
In [23]: def precision_at_k(row, k=10):
          return len(set(row['content_based_recommended']) & set(df_answer[row['personId']]))

          evaluation_result = df_user_preference.apply(precision_at_k, axis=1)
          print("Average Precision At K:", np.average(evaluation_result))
          # Average Precision At K: 0.5700877192982456

Average Precision At K: 0.5700877192982456
```